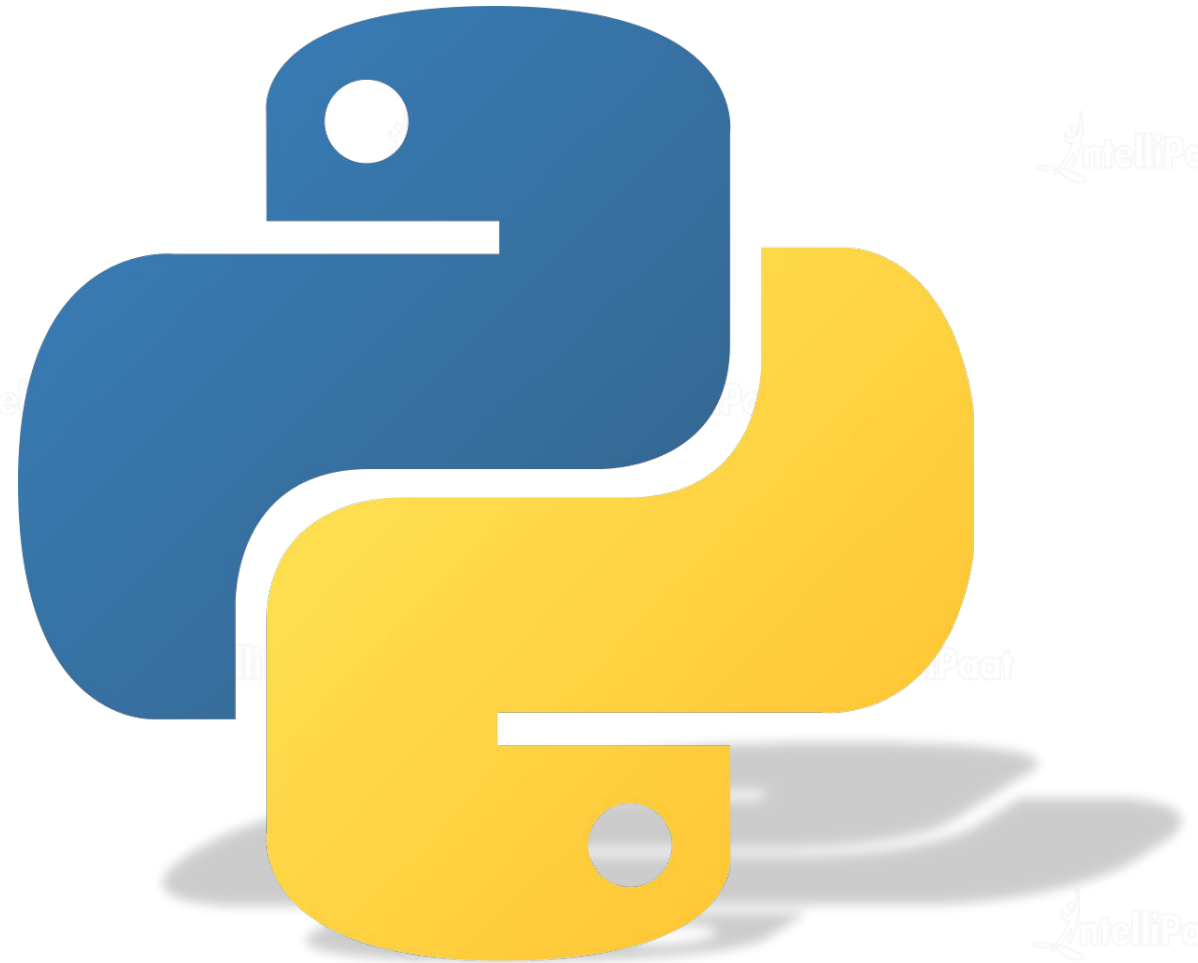




Time Series Forecasting



Agenda

01

Loading Time Series Data

02

Time Series Visualization

03

Time Series Analysis

04

Stationarity in Time Series

05

Time Series Differencing

06

Removing trends and
Seasonality

07

Pyramid ARIMA

08

ARIMA Model for Time Series
Forecasting

09

SARIMA Model for Times Series
Forecasting

10

Inferences

Problem Statement

Problem Statement

The sudden influx of passengers can be an overwhelming task for the airport authorities across the globe. With time series data of the passengers over the years, one can forecast the estimated number of passengers that will fly for the specific months of the year. Use the Passenger Dataset to forecast the number of passengers for the upcoming years.



The Dataset contains the monthly passenger data from **January 1949** to **December 1960**. A decade long monthly information of the data converted in a time series.

The data contains two columns – **Month** and **#Passengers**.



Loading Time Series Data

How to load the Time Series Data?

We will make use of the Pandas library in Python Programming to load our dataset. We will follow the following steps to load the time series data.

01

We will use the **read_csv()** method to import the dataset from the local machine.

02

After importing the dataset, we will convert the Month column to datetime, using the **to_datetime()** method from the pandas module.

03

Since we have two columns, we will convert the month column to index for easier computation.



```
import pandas as pd

#importing the dataset from local machine
data = pd.read_csv("AirPassengers.csv")
#converting the data to datetime object type
data['Month'] = pd.to_datetime(data['Month'])
#changing the column to index
data.index = data['Month']
del data['Month']
#displaying the data
data.head()
```

#Passengers	
Month	
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

Time Series Visualization

Plotting the Time Series

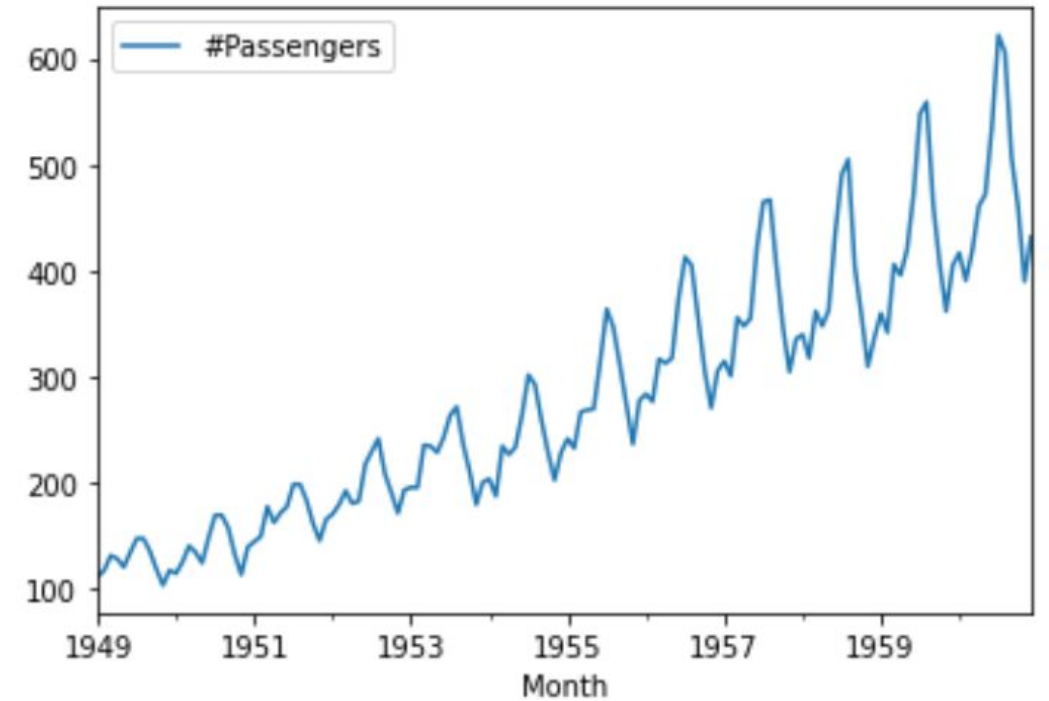
Plotting the Time series can help in analyzing the visible trends, seasonality in the data. We can use matplotlib or seaborn library to create to create simple line plots to study the data.



Hands on

```
import matplotlib.pyplot as plt  
import seaborn as sns
```

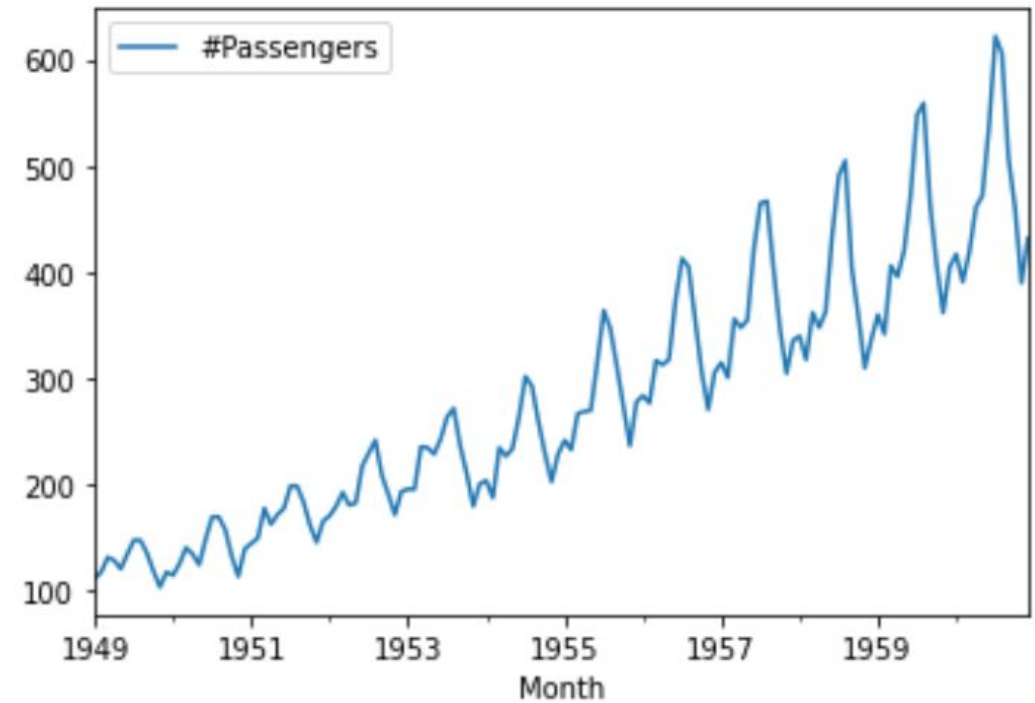
```
data.plot()
```



Time Series Analysis

Analyze the Time Series

From the plot drawn in the previous section, we can clearly see a few patterns. There is an upward trend in the data with seasonality. To understand this, let's take a look at various patterns to identify from time series plots.



A trend is a pattern when there is a long time increase or decrease in the data.

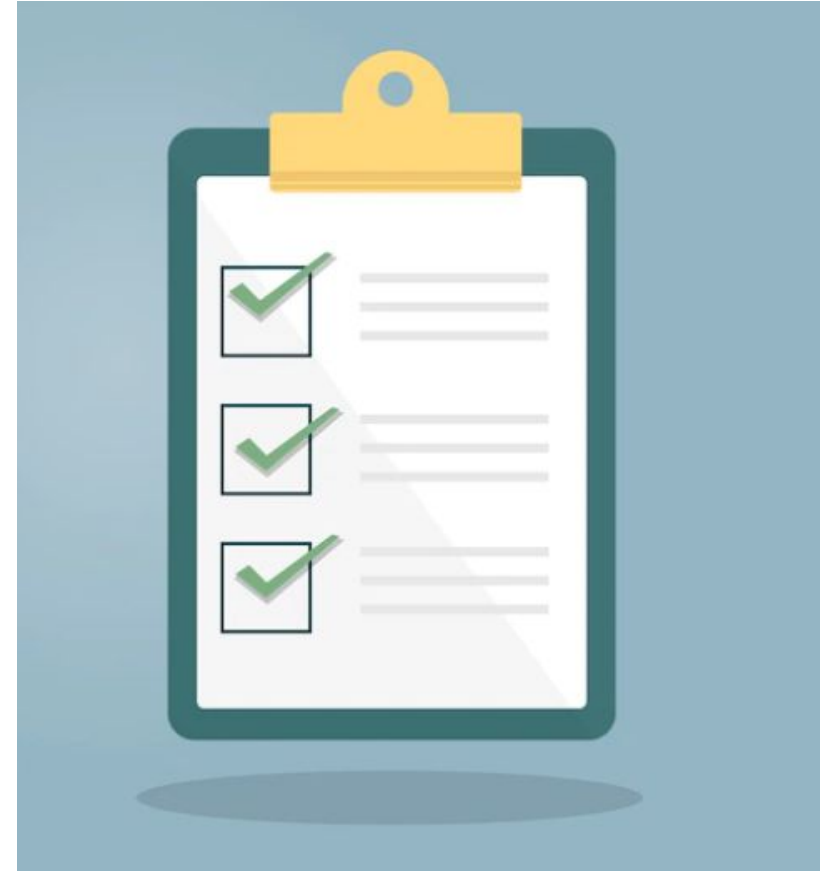
Seasonality pattern happens in equal intervals based on parameters like week, month or a year.

Cyclic pattern showcases the rise and fall in the data that are not fixed with respect to a specific frequency like in a seasonal pattern.

Stationarity in a Time Series

What is Stationarity?

Before we can begin modeling with the ARIMA model for forecasting, we have to make sure the time series data is stationary. In simple terms, if a data consists of trends and seasonality, the data is not going to be a stationary data. In our case, we have both, so let's take a look at how we can fix this.



How to Find the Stationarity of a Time Series?

To find the stationarity in a data, we can use the statistical test such as Augmented Dickey Fuller test.

Here, we will take two hypotheses, a null hypotheses and an alternate hypotheses. If we are not able to reject the null hypotheses after the computation, the time series is a stationary series.

To reject the null hypotheses, the following must be true:

1. If the p-value after the adfuller test is greater than 0.05, we fail to reject the hypotheses.
2. If the p-value is less than 0.05, we can reject the null hypotheses and assume that the time series is stationary.

In Python, we can make use of the **statsmodels.tsa.stattools** package that provides the **adfuller** module to conduct the augmented dickey-fuller test on the time series.

```
#checking the stationarity of the data
from statsmodels.tsa.stattools import adfuller

result = adfuller(data['#Passengers'])
print(result)
```

The p-value from the result is on the index – 1, and we can check if the data is stationary or not. In our case, we have the p-value more than 0.05 and thus, we cannot reject the null hypotheses and assume the data to be non-stationary.

Time Series Differencing

Differencing in time series is the process of reducing the non-stationary time series to a stationary time series with a series of subtraction operations i.e. subtracting the observations from one another.

$$\text{Diff}(t) = x(t) - x(t-1),$$

Where $\text{Diff}(t)$ is the differenced series, $x(t)$ is the observation at given time t , and the $x(t-1)$ is the subsequent observation.

Other ways to remove Stationarity

There are several ways to remove stationarity from the time series you can choose from.

Differencing is the most common technique to remove stationarity.

You can use power transformation for removing the stationarity.

Log transformations of the time series is another technique to remove the stationarity.

Removing Trends and Seasonality

Removing Trends

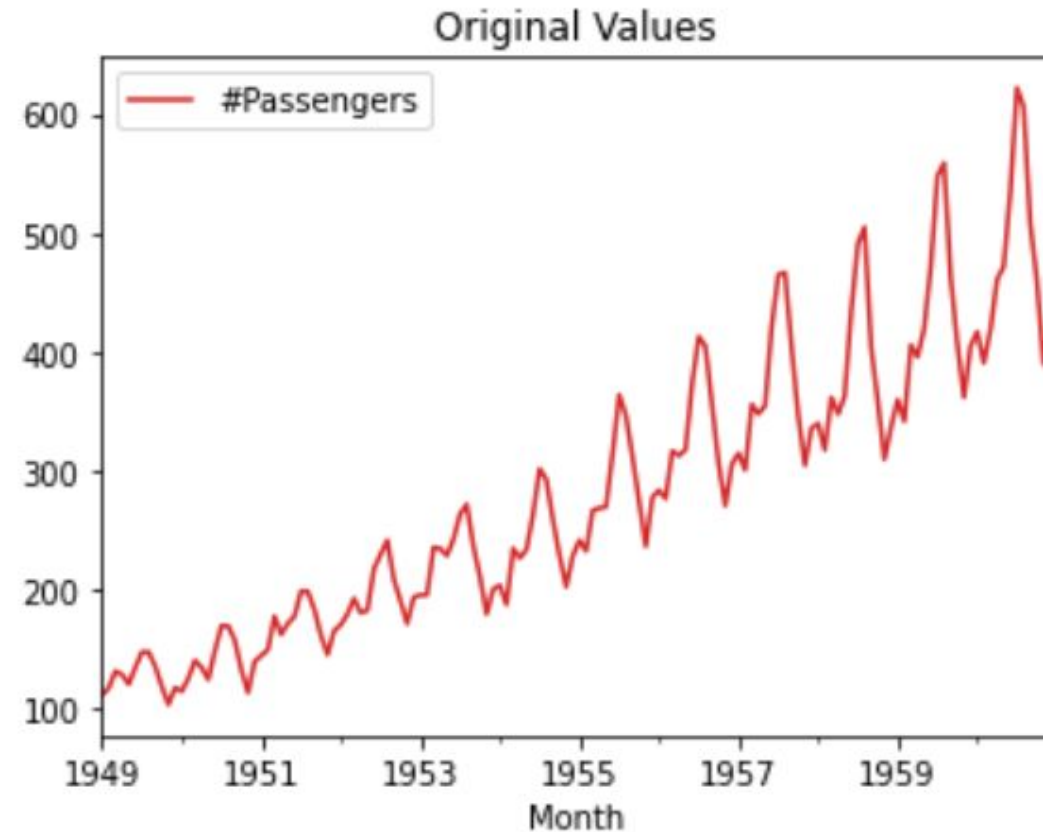
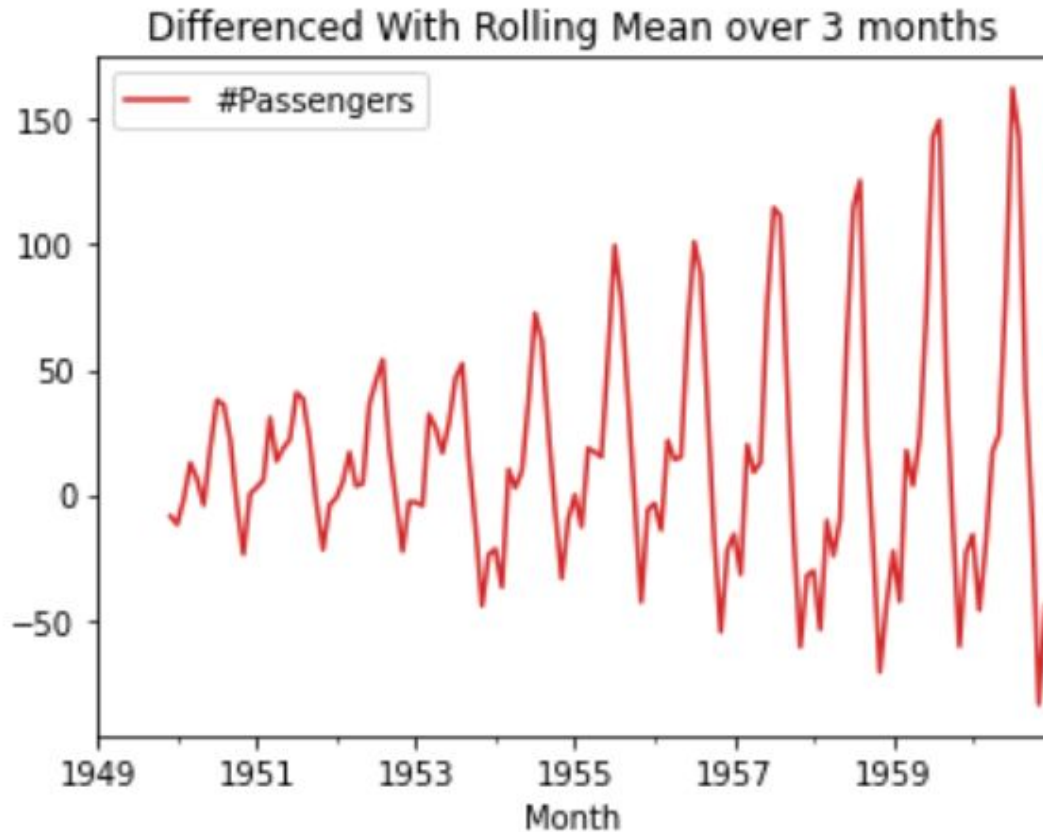
we will remove the trend from the time series using the rolling mean differencing as shown in the example below.

```
#removing trend using the rolling mean differencing
rolling_mean = data.rolling(window=12).mean()
rolling_mean_detrended = data - rolling_mean

ax1 = plt.subplot(121)
rolling_mean_detrended.plot(figsize=(12,4),color="tab:red",
                             title="Differenced With Rolling Mean over 12 months",
                             ax=ax1)

ax2 = plt.subplot(122)
data.plot(figsize=(12,4),
           color="tab:red",
           title="Original Values",
           ax=ax2)
```

Removing Trends



Time Series Decomposition

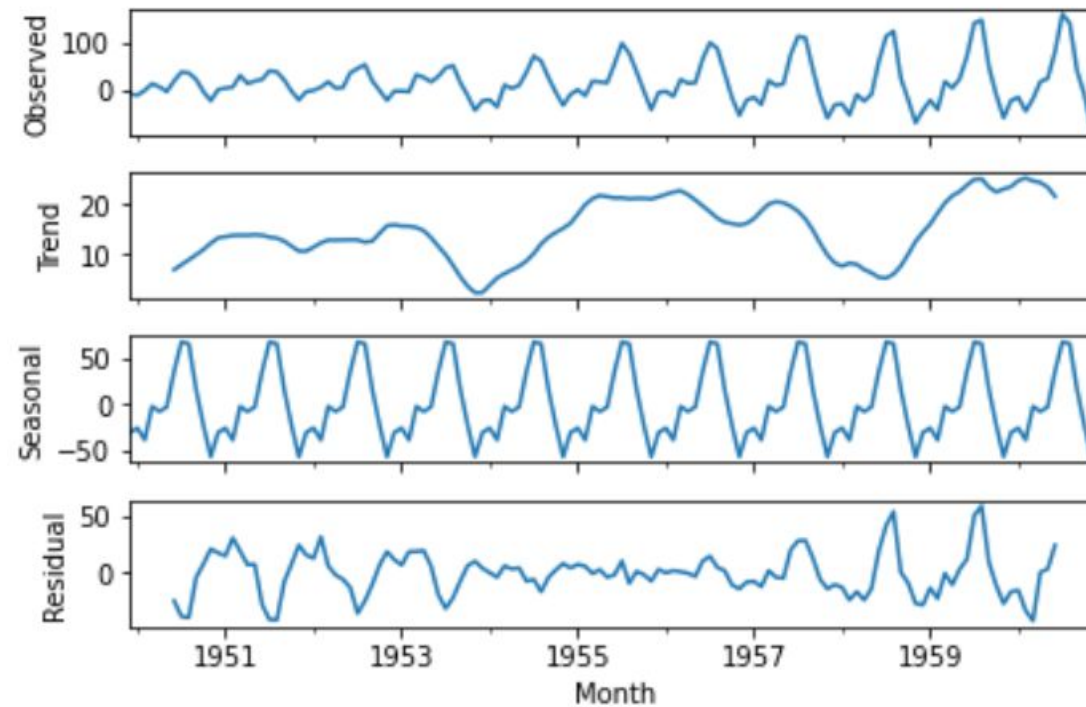
We can get the decomposition of the time series data using the `seasonal_decompose()` method from the `statsmodels.tsa.seasonal`.

```
from statsmodels.tsa.seasonal import seasonal_decompose
decompose_result = seasonal_decompose(rolling_mean_detrended.dropna())

decompose_result.plot()
```

Time Series Decomposition

We can check the trend and seasonality using the `seasonal_decompose()` method.



Removing Seasonality

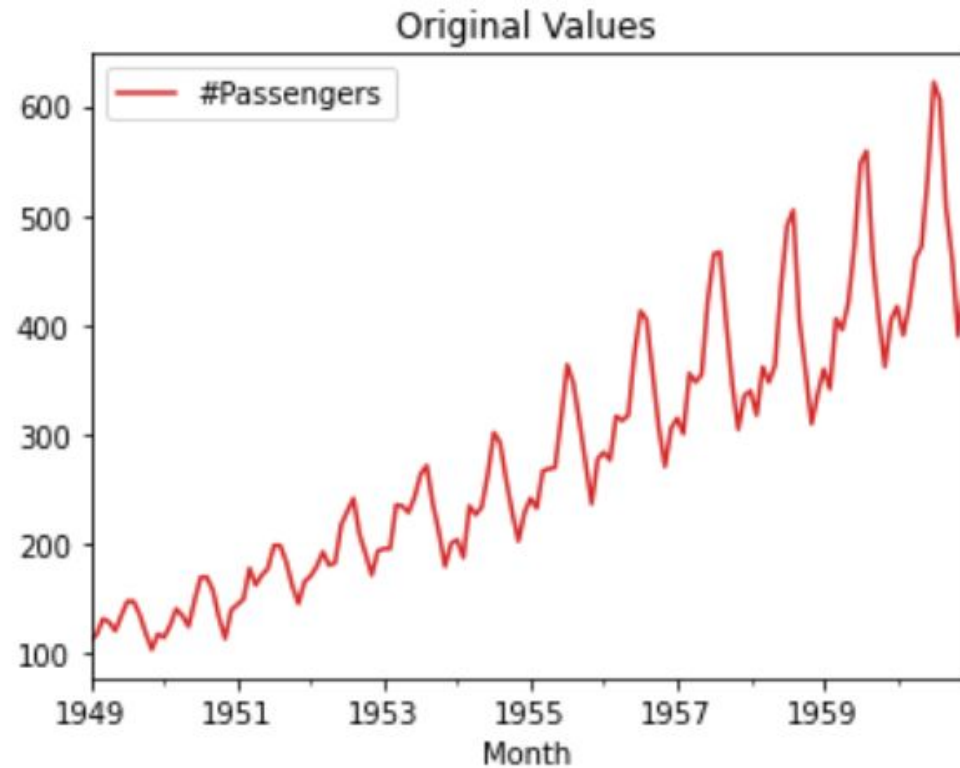
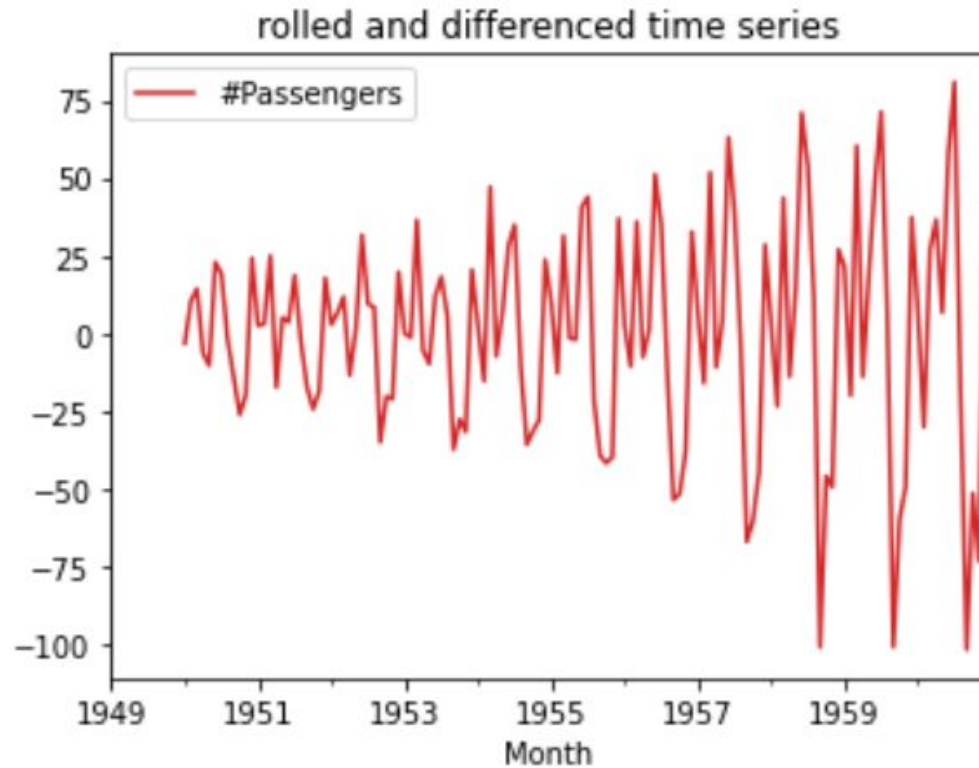
we will remove the seasonality from the time series using the rolling mean differencing as shown in the example below.

```
#removing seasonality from the time series
rolling_mean_detrended_diff = rolling_mean_detrended - rolling_mean_detrended.shift()

ax1 = plt.subplot(121)
rolling_mean_detrended_diff.plot(figsize=(12,4),
                                color="tab:red",
                                title="rolled and differenced time series",
                                ax=ax1)

ax2 = plt.subplot(122)
data.plot(figsize=(12,4),
           color="tab:red",
           title="Original Values",
           ax=ax2)
```

Removing Seasonality



Adfuller test on the new detrended time series to check the stationarity of the time series.

```
result = adfuller(rolling_mean_detrended_diff['#Passengers'].dropna())  
print(result)
```

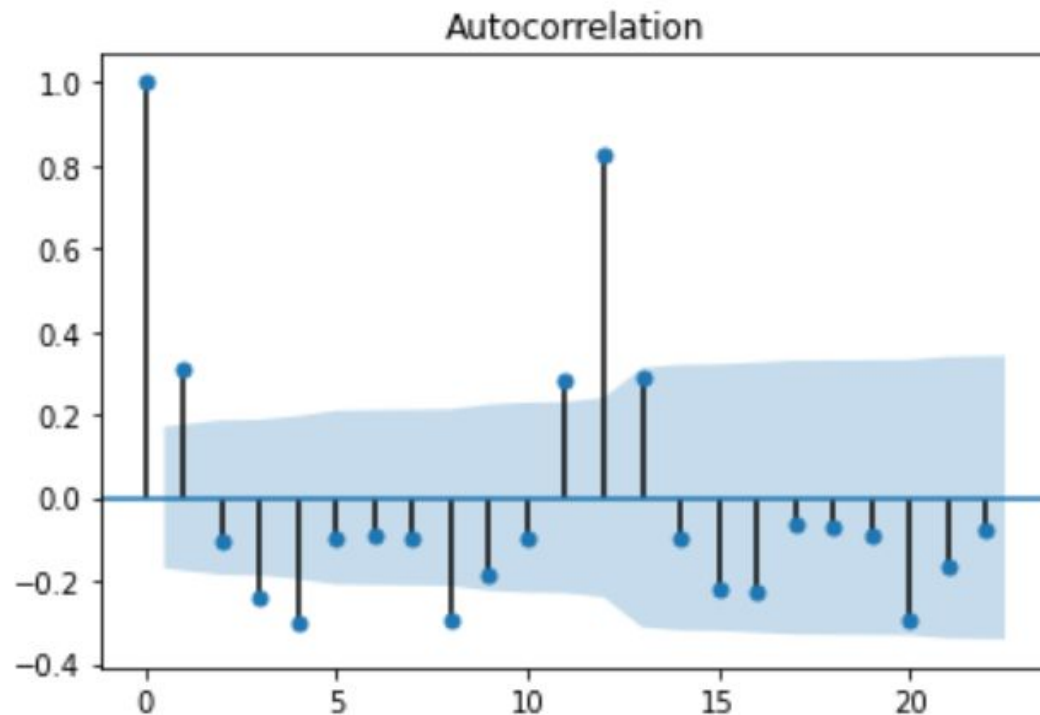
The p-value from the result is less than 0.05, therefore we can reject the null hypotheses and consider the time series to be stationary.

Autocorrelation & Partial Autocorrelation

Autocorrelation

```
from pandas.plotting import autocorrelation_plot
from statsmodels.graphics.tsaplots import plot_acf

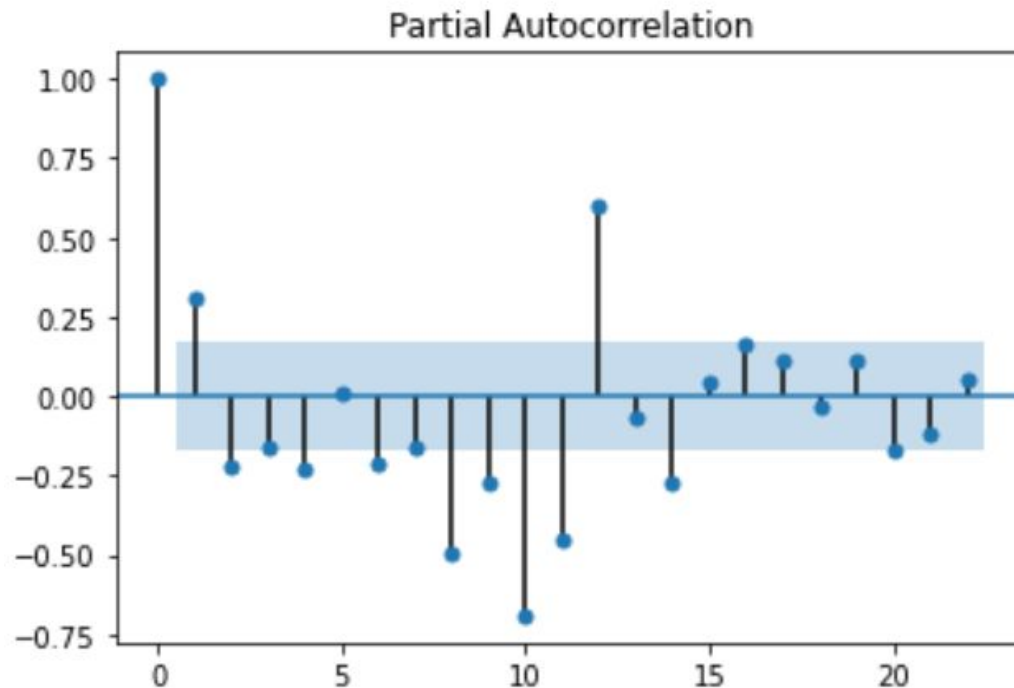
plot_acf(rolling_mean_detrended_diff['#Passengers'])
```



Autocorrelation and partial autocorrelation plots can be used to determine the ARIMA orders for the AR and MA models.

Partial Autocorrelation

```
from statsmodels.graphics.tsaplots import plot_pacf  
  
plot_pacf(rolling_mean_detrended_diff['#Passengers'])
```



Pyramid ARIMA

What is Pyramid ARIMA?

The pyramid ARIMA or pmdarima package from python programming can be used to compute the order of the ARIMA model using the auto_arima module.

```
!pip install pmdarima
from pmdarima import auto_arima
|
order = auto_arima(rolling_mean_detrended_diff['#Passengers'], trace=True)
order.summary()
```

The auto_arima model searches the best order for the model that has the lowest AIC value.

Auto ARIMA in Pyramid ARIMA

```
ARIMA(1,0,0)(0,0,0)[0] : AIC=inf, Time=0.17 sec
ARIMA(1,0,2)(0,0,0)[0] : AIC=inf, Time=0.18 sec
ARIMA(0,0,2)(0,0,0)[0] : AIC=1295.960, Time=0.05 sec
ARIMA(2,0,0)(0,0,0)[0] : AIC=1296.291, Time=0.04 sec
ARIMA(2,0,2)(0,0,0)[0] : AIC=1242.347, Time=0.16 sec
ARIMA(3,0,2)(0,0,0)[0] : AIC=inf, Time=0.37 sec
ARIMA(2,0,3)(0,0,0)[0] : AIC=inf, Time=0.37 sec
ARIMA(1,0,3)(0,0,0)[0] : AIC=inf, Time=0.25 sec
ARIMA(3,0,1)(0,0,0)[0] : AIC=inf, Time=0.14 sec
ARIMA(3,0,3)(0,0,0)[0] : AIC=1235.785, Time=0.47 sec
ARIMA(4,0,3)(0,0,0)[0] : AIC=inf, Time=0.40 sec
ARIMA(3,0,4)(0,0,0)[0] : AIC=inf, Time=0.53 sec
ARIMA(2,0,4)(0,0,0)[0] : AIC=inf, Time=0.95 sec
ARIMA(4,0,2)(0,0,0)[0] : AIC=inf, Time=0.55 sec
ARIMA(4,0,4)(0,0,0)[0] : AIC=inf, Time=1.28 sec
ARIMA(3,0,3)(0,0,0)[0] intercept : AIC=1236.433, Time=2.16 sec
```

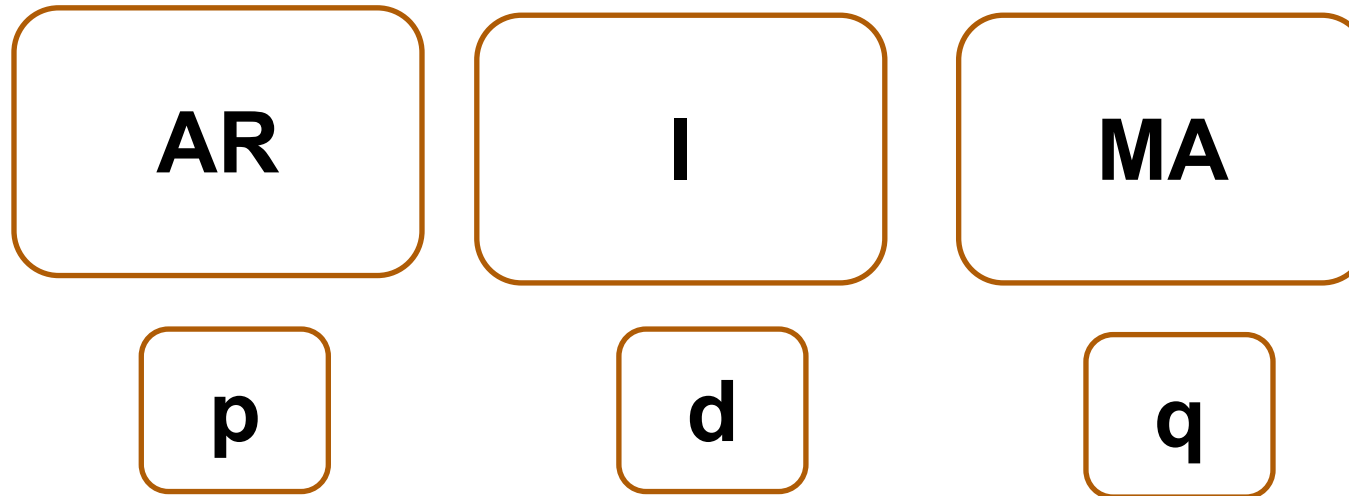
Best model: ARIMA(3,0,3)(0,0,0)[0]
Total fit time: 9.725 seconds

From the output, the best order for the ARIMA model is (3,0,3), for (p,d,q) respectively.

ARIMA Model For Time Series Forecasting

What is ARIMA?

ARIMA model is the combination of Autoregressive(AR), Integrated (I), and Moving Average(MA) models.



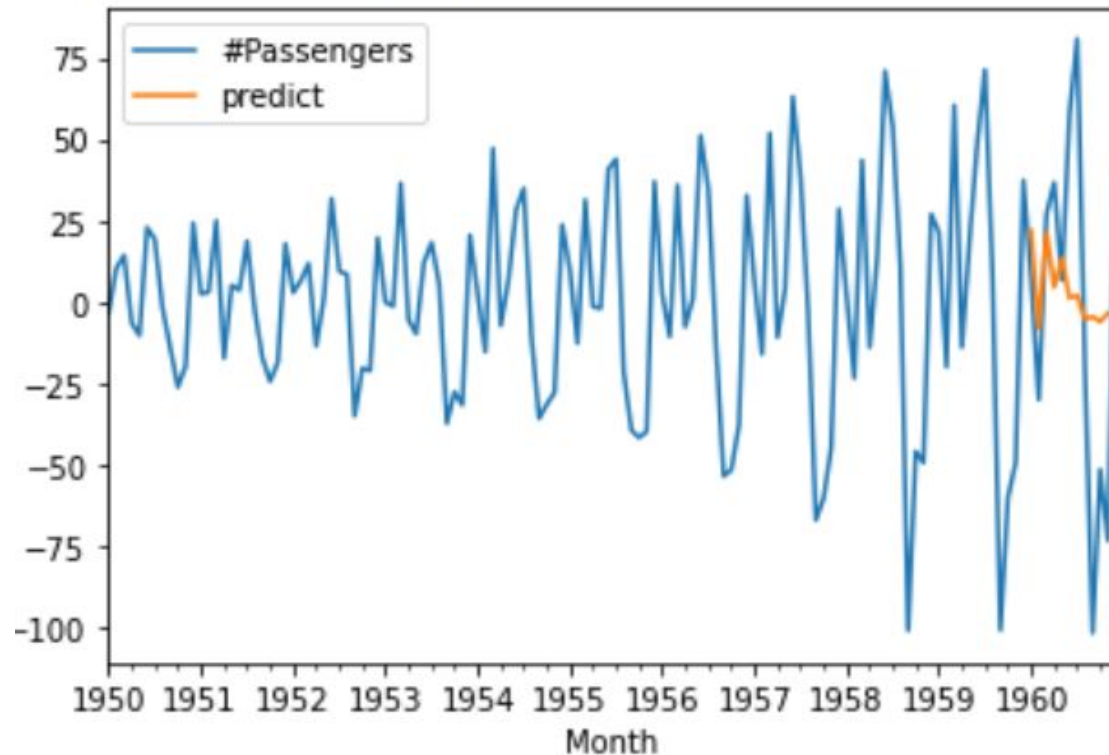
Here, p, d and q are the order of AR, order of differencing and order of MA respectively. We have already calculated these values using the `auto_arima`, or we can deduce these values using the ACF and PCF.

```
from statsmodels.tsa.arima_model import ARIMA

train = rolling_mean_detrended_diff.iloc[:120]['#Passengers']
test = rolling_mean_detrended_diff.iloc[121:]['#Passengers']

model = ARIMA(train, order=(3,0,3))
model_fit = model.fit()
model_fit.summary()

rolling_mean_detrended_diff['predict'] = model_fit.predict(start= len(train),
                                                           end=len(train)+len(test)- 1,
                                                           dynamic=True)
rolling_mean_detrended_diff[['#Passengers','predict']].plot()
```



As you can see, the predictions are way off the actual values from the test set. Therefore, we can move to the seasonal ARIMA model for our forecasting.

SARIMA Model For Time Series Forecasting

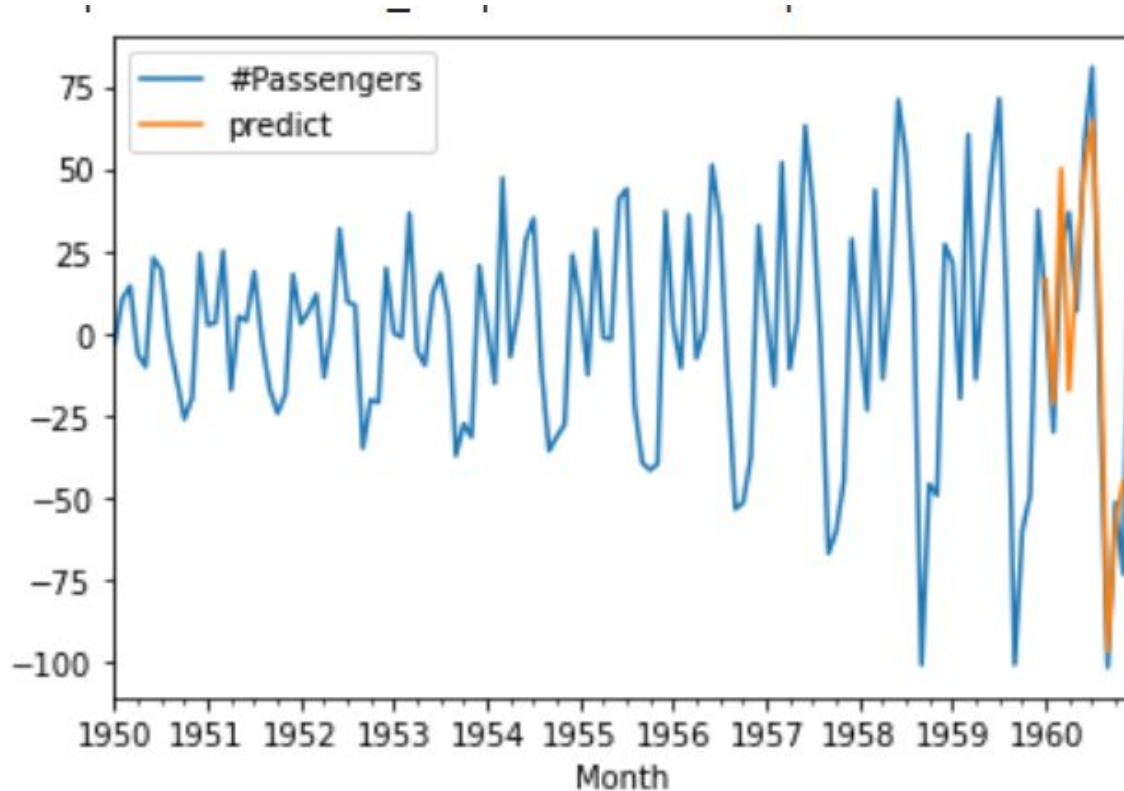
What is Seasonal ARIMA?

In the seasonal ARIMA model, we have to specify the seasonal order as well. The seasonal order remains the same as the ARIMA order, and we can add the periodic order in the seasonal order according to the periodicity.

```
from statsmodels.tsa.statespace.sarimax import SARIMAX

model = SARIMAX(train, order=(3,0,3), seasonal_order=(3,0,3,12))
model = model.fit()
```

```
rolling_mean_detrended_diff['predict'] = model.predict(start=len(train) ,
                                                         end=len(train)+len(test)- 1,
                                                         dynamic=True)
rolling_mean_detrended_diff[['#Passengers', 'predict']].plot()
```

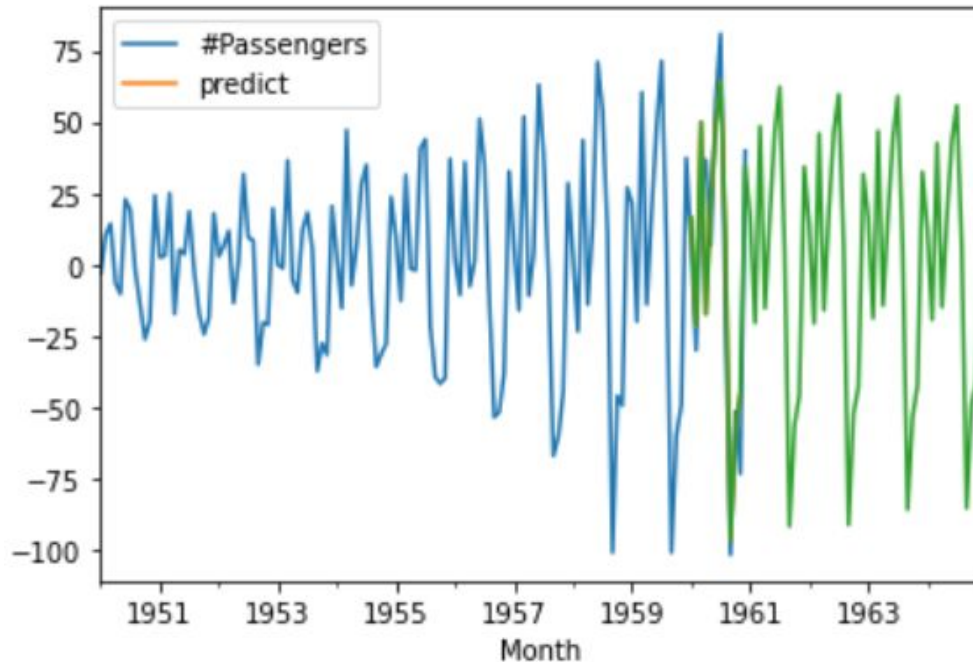


Here, we can see the predicted values on the test set are more accurate than the ARIMA model. Therefore we have successfully created a Time series forecast model. Now we will use this model to forecast the time series.

Inferences

```
#predicting the projections for the next 5 years
```

```
forecast = model.forecast(steps=60)  
rolling_mean_detrended_diff.plot()  
forecast.plot()
```



We had trained the model on the rolling_mean_detrended_diff values, therefore the predictions are aligned to the same. We can train the model with the original dataset, and add the order of differencing manually and get the predictions on the actual values.



India: +91-7847955955

US: 1-800-216-8930 (TOLL FREE)



support@intellipaate.com



24/7 Chat with Our Course Advisor